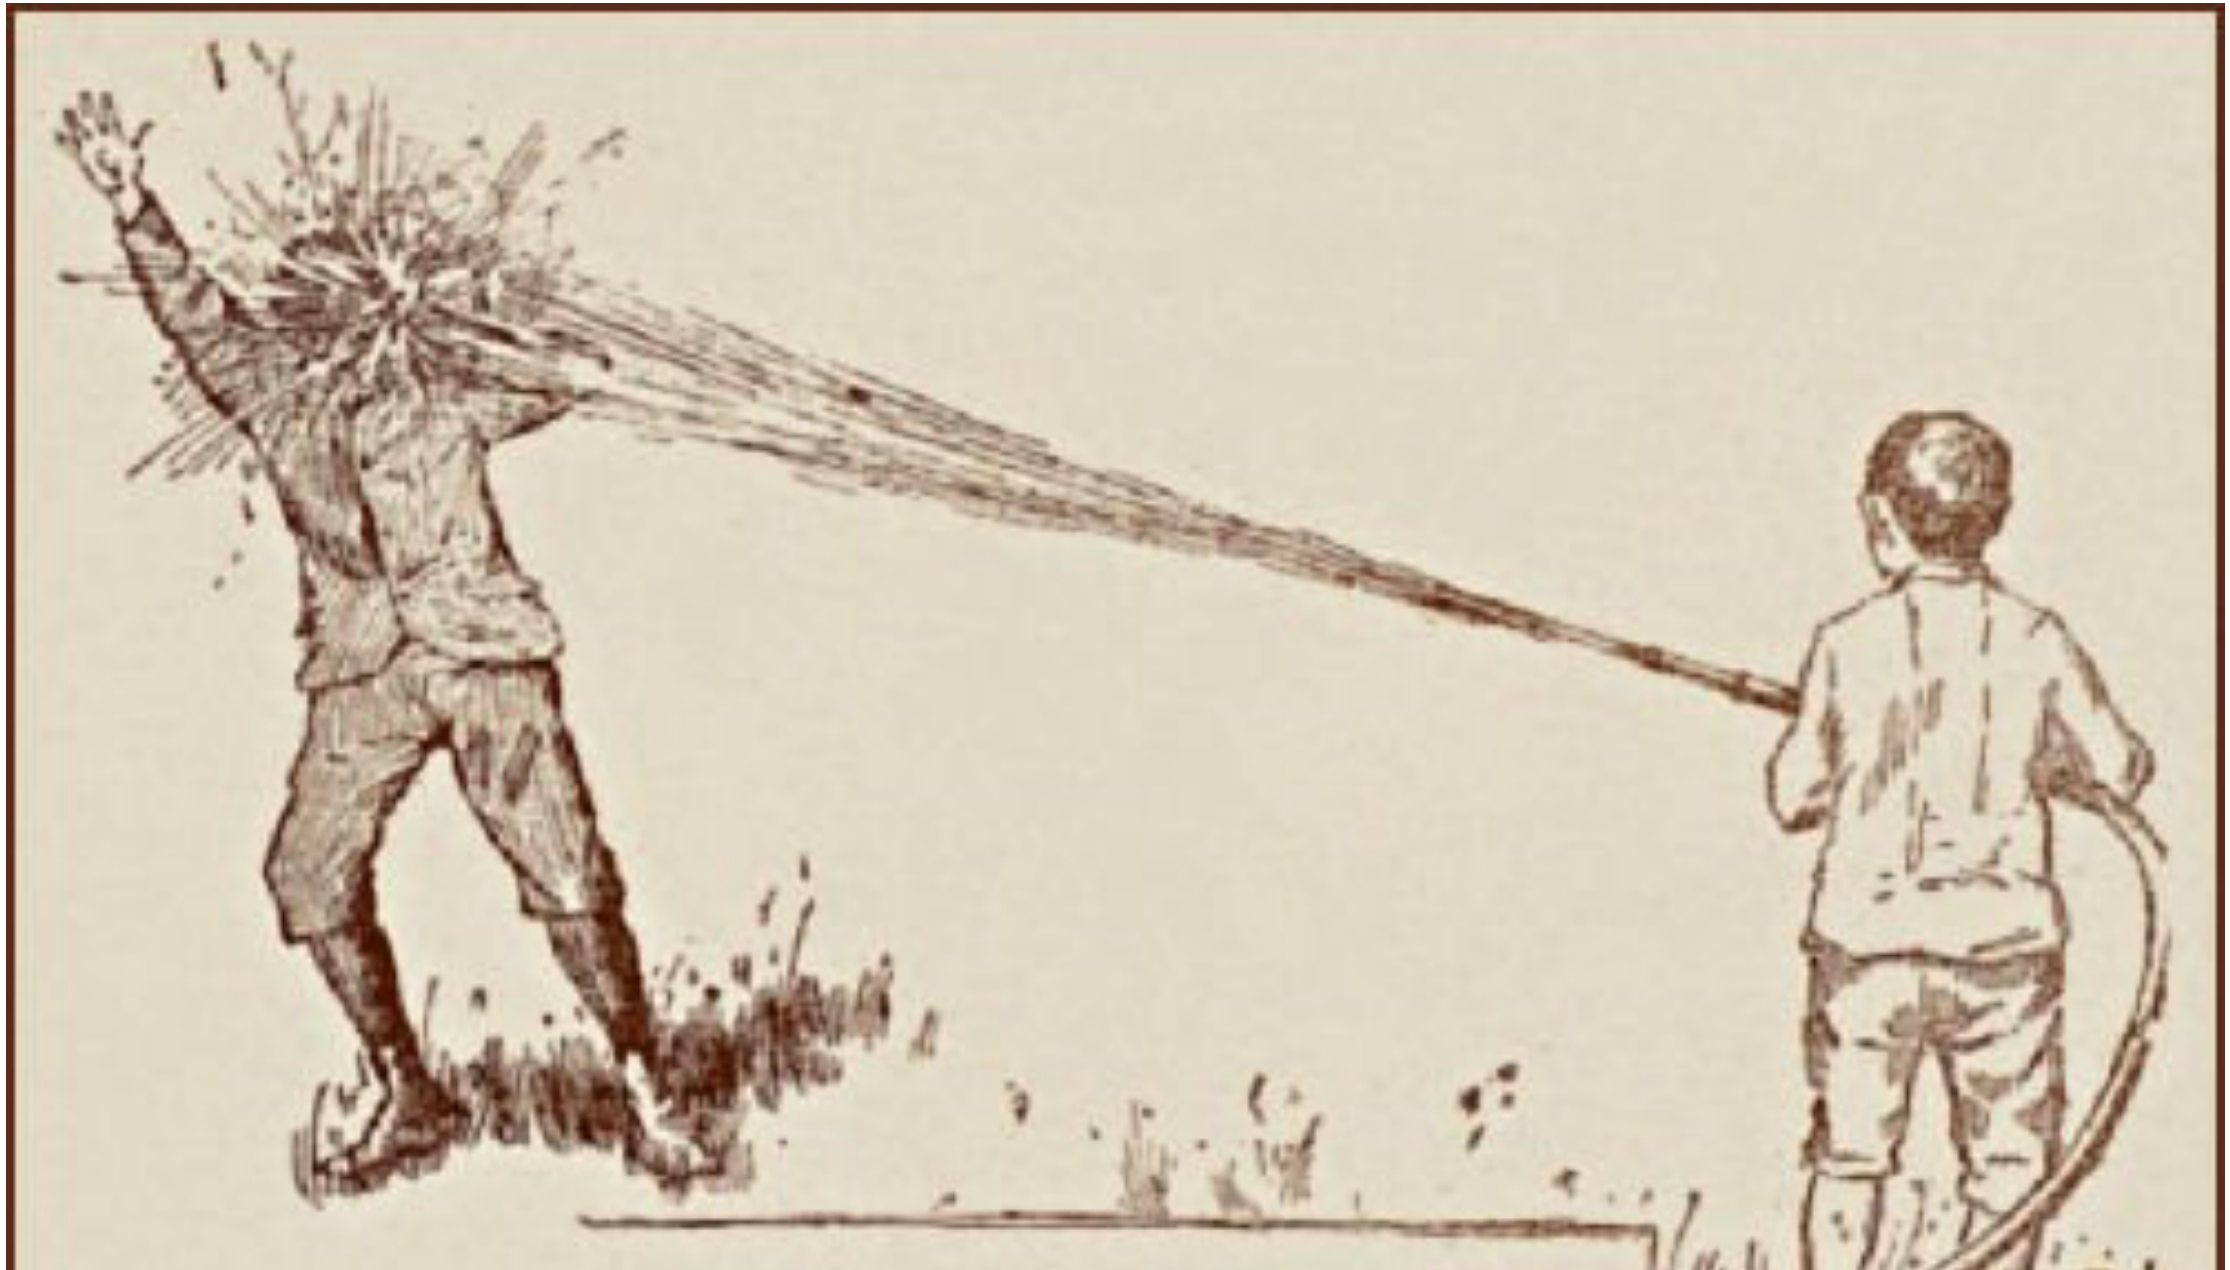


Let's Drink from the Firehose




Twitter and MongoDB

MongoDB


- A Database Server
- Contains Databases that contain Collections that contain Documents
- Documents are stored using a JSON-like format; play very well with `dict` objects

MongoDB



Blueprint

- Will use default images from Docker Hub for both services
- We will run them on two different systems
 - Jupyter - either local or on AWS
 - MongoDB - on a separate AWS server



NoSQL

- MongoDB is a NoSQL database
- No schema required
- Just add a document (dict) to a collection

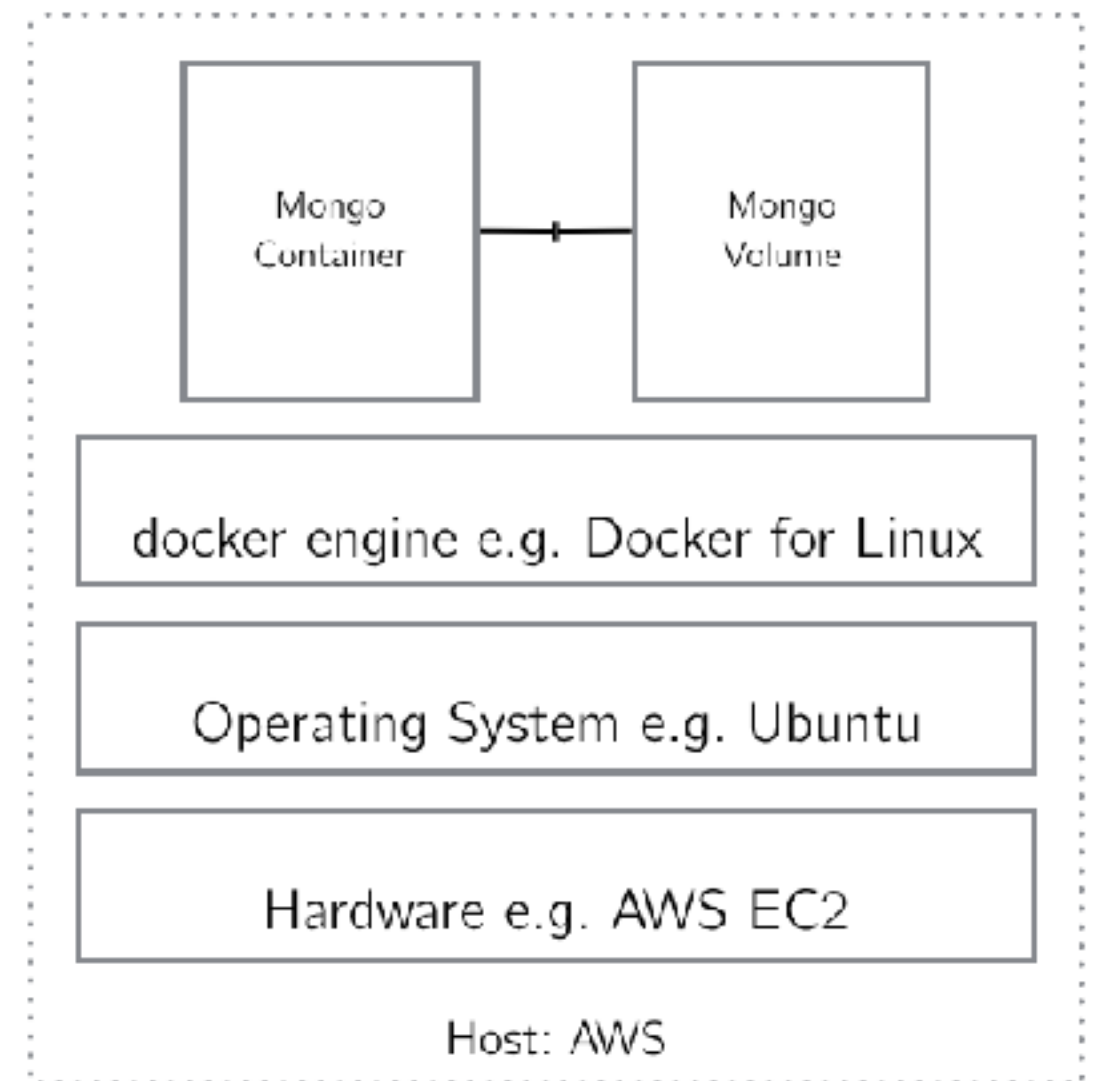
NoSQL



You == Han Solo
Your data == Chewbacca

Configuring Mongo on AWS

- Considerations:
 - Networking
 - Solve via AWS Security Groups
 - Data Persistence
 - Solve via Docker Volumes



Configuring Mongo on AWS

- Note that this configuration will use Mongo on a separate instance from the instance on which you are running Jupyter.
- If you are running Jupyter on AWS, you will need a second `t2.micro`.
- Accessing a database managed by Docker from a different AWS instance is actually easier than accessing a database managed by Docker on the same system.

Set up a new t2.micro

- From the AWS EC2 Dashboard, select “Launch Instance.”
- On the Choose AMI tab, choose Ubuntu Server 16.04.
- On the Choose Instance Type tab, choose t2.micro.
- On the Add Storage tab, use the default setting of 8GB.
- On the Configure Security Group tab, choose “Create a new security group.”
- a. Confirm that inbound SSH traffic can be accepted over port 22 from anywhere.
- b. Add a rule that accepts inbound traffic over port 2376 from anywhere. This port will allow you to pull images from Docker Hub.
- c. Add a rule that accepts inbound traffic over port 27016 from anywhere. This is the default port for accessing MongoDB.
- Review and launch an instance, taking care to confirm that you have access to the SSH keys stored with your AWS account.

Configure the new t2.micro

- Take note of the IP address of the newly configured AWS instance.
- SSH into the instance using that IP address.
- Install Docker via a shell script.
- Add the ubuntu user to the docker group.
- Log out and back in.

```
(local) $ ssh ubuntu@255.255.255.255  
(remote) $ curl -sSL https://get.docker.com/ | sh  
(remote) $ sudo usermod -aG docker ubuntu
```

Run Mongo via Docker

- Pull the mongo image

```
$ docker pull mongo
```

- Create a New Data Volume

```
$ docker volume create mongo-dbstore
```

- Launch MongoDB as a Persistent Service

```
$ docker run -d --name this_mongo \  
    -v mongo-dbstore:/data/db \  
    -p 27017:27017 mongo
```

Verify MongoDB Installation

- You can verify that you are running the mongo service by connecting to the running MongoDB via the MongoDB client, `mongo`, issued via `docker exec`.
- To do this, connect and then insert a trivial document to a mongo collection. You are inserting the JSON object `{"foo": 1}` into the collection `test`. You then search for the document you inserted using the `.find()` command.

```
$ docker exec -it this_mongo mongo
> db.test.insert({"foo":1})
> db.test.find()
```

Using MongoDB with Jupyter

- You will need to install the necessary Python library, `pymongo`

```
!pip install pymongo
```

- This should be run from a Jupyter server that is not on the same AWS instance as your Mongo server.

pymongo

- pymongo is a Python module containing the MongoDB tools recommended for working with the database.
- You begin by instantiating a connection to MongoDB using `pymongo.MongoClient`.
- Here, you use the IP address of your AWS instance on which MongoDB is running.

```
from pymongo import MongoClient  
client = MongoClient('255.255.255.255', 27016)
```

pymongo

- pymongo has a very useful “get or create” mechanism for both databases and collections.

```
client.database_names()
```


pymongo

- Databases and collections are accessed using either attribute-style (`client.database_name`) or dictionary-style (`client['test-database']`).
- If the database exists, this method will return a reference to the existing database or collection (“get”). If the database does not exist, this method will create the database or collection and then return a reference to it (“create”).

pymongo

- The creation happens at the time of insertion of a document.

```
db_ref = client.my_database  
client.database_names()
```

```
coll_ref = db_ref.my_collection  
client.database_names(), db_ref.collection_names()
```

```
sample_doc = {"name": "Joshua", "message": "Hi!",  
'my_array' : [1,2,3,4,5,6,7,9]}  
coll_ref.insert_one(sample_doc)
```

```
client.database_names(), db_ref.collection_names()
```


Mongo and Twitter

- To demonstrate a simple usage for MongoDB with Jupyter, you will implement a basic Twitter streamer that inserts captured tweets into a MongoDB collection.
- Twitter data represents an ideal use case for the NoSQL MongoDB.
- Each tweet obtained via the Twitter API is received as an unstructured nested JSON object.

Mongo and Twitter


- Adding such an object to a SQL database would be a non-trivial task by any measure involving numerous foreign keys and JoinTables as the user seeks to manage each of the one-to-one, one-to-many, and many-to-one relationships built into the tweet.
- Adding such an object to Mongo, on the other hand, is a trivial task.
- MongoDB's native Binary JSON (BSON) format was designed precisely to accept such an object.

Let's Drink from the Firehose



Obtain Twitter Credentials

- In order to follow along, you must obtain API credentials for accessing the Twitter API.
This is done by creating a Twitter application.
- In order to do this, follow these steps:
 1. Visit <https://apps.twitter.com> and sign in.
 2. Choose “Create New App”.
 3. Give the new app a name, description, and website. For your purposes, the values of these responses are irrelevant, although the website will need to have a valid URL structure.
 4. Agree to the Developer Agreement and click “Create your Twitter Application”.



Obtain Twitter Credentials

- Next, you will need to access your credentials on the “Keys and Access Tokens” tab.

- You will need a total of four values:
 1. A consumer key (API Key)
 2. A consumer secret (API Secret)
 3. An access token
 4. An access token secret






Test_CJP

[Details](#) [Settings](#) [Keys and Access Tokens](#) [Permissions](#)

Test OAuth

Application Settings

Keep the “Consumer Secret” a secret. This key should never be human readable in your application.

Consumer Key (API Key)		
Consumer Secret (API Secret)		
Access Level	Read and write (modify api permissions)	
Owner	joshuacook	
Owner ID		

Your Access Token

This access token can be used to make API requests on your own account's behalf. Do not share your access token secret with anyone.

Access Token		
Access Token Secret		
Access Level	Read and write	
Owner	joshuacook	
Owner ID		

Load Twitter Credentials

- Load Twitter Credentials as Strings
- Replace this with your credentials:

```
CONSUMER_KEY = None  
CONSUMER_SECRET = None  
ACCESS_TOKEN = None  
ACCESS_SECRET = None
```


Install the `twitter` library

- I prefer the `twitter` library over `tweepy`. I've found it to be better for streaming. Others have found `tweeps` better for historical data.

```
!pip install twitter
```

Authentication

- You next instantiate a `twitter.OAuth` object using the Python twitter module and the credentials you have just loaded.
- You will use this object to facilitate your connection to Twitter's API.


```
from twitter import OAuth
oauth = OAuth(ACCESS_TOKEN, ACCESS_SECRET,
             CONSUMER_KEY, CONSUMER_SECRET)
```

Collect Tweets by Geolocation

- For this example, you will be using Twitter's Public Stream.
- Applications that are able to connect to a streaming endpoint will receive a sample of public data flowing through Twitter and will be able to do so without polling or concern of API rate limits.
- In other words, the Public Stream is a safe and sanctioned way to collect a sample of live public tweets.
- That said, even this sample will return a great deal of unordered data.

Collect Tweets by Geolocation

- In order to provide a modicum of order to your Twitter stream, you will restrict incoming tweets using a geolocation bounding box, or bbox. You can easily obtain a bbox for a location of interest using the Klokantech BoundingBox Tool.
- Let's obtain a bbox for Santa Monica, California in the United States, making sure to select CSV Raw as the copy and paste format.



```
los_angeles_bbox = "-118.551346,33.96666,-118.443428,34.05056"
```

Instantiate a TwitterStream

- Finally, you instantiate a `twitter.TwitterStream` object you will use to collect tweets.
- `twitter.TwitterStream` provides an interface to the Twitter Stream API in Python.
- The result of calling a method on this object is an iterator that yields tweets decoded from the Twitter stream as JSON objects.

```
from twitter import TwitterStream
```

```
twitter_stream = TwitterStream(auth=oauth)
```

```
twitterator =
```

```
twitter_stream.statuses.filter(locations=los_angeles_bbox)
```


Insert Tweets Into Mongo

- Twitter is a wonderful source of messy, “real” data.
- Wrangling it into a database is where MongoDB truly shines.
- Using your `twitterator` object and the `.insert_one()` class function this can be done in a single line of code.

```
coll_ref.insert_one(next(twitterator))  
coll_ref.count()  
coll_ref.find_one()
```


Directed Acyclic Graph

- A **graph** is a collection of nodes and edges
- A **directed graph** is a graph where each edge is directional.
- A **directed acyclic graph** (DAG) is a directed graph with no cycles.
- Once we have passed through a node, we never return to it.




Directed Acyclic Graph

- The DAG is a great way to model a programmed process.



Directed Acyclic Graph

- The DAG for our Tweet Streamer is



Simple Parallelism

- What if I wanted to do more than one streamer at a time?
- Trivially, this can be done by running n notebooks, where n is the number of cpus on your computer.

